

# Towards behavioral consistency in multi-modeling

Tim Kräuter\*, Harald König\*<sup>†</sup>, Adrian Rutle\*, Yngve Lamo\*, and Patrick Stünkel<sup>‡</sup>

\*Western Norway University of Applied Sciences, Bergen, Norway

<sup>†</sup>University of Applied Sciences, FHDW, Hannover, Germany

<sup>‡</sup>Haukeland Universitetssykehus, Bergen, Norway

**ABSTRACT** Multiple heterogeneous interacting systems are needed to realize the requirements of complex domains. Describing the interactions between these systems and checking their global behavioral consistency is a general, well-known challenge in software engineering. To address this challenge, model-driven software engineering utilizes abstract representations of the constituting systems and their interactions, resulting in a *multi-model* representing the overall system. In such a multi-modeling setting, global consistency rules must be satisfied by a set of heterogeneously typed models to guarantee a desired *global behavior*. In this paper, we propose a novel approach for behavioral consistency management of heterogeneous multi-models. The approach introduces a workflow in which we (i) define which behavioral models in the multi-model *can interact*, (ii) specify consistency requirements as *global behavioral properties*, (iii) align the individual models by specifying *how they interact*, (iv) generate a formal specification of the *global behavior*, and finally, (v) check the *global behavioral properties*, which should be satisfied by the multi-model. Our approach is independent of the particular formalism used in the generated specification, and we currently support graph transformations (Groove) and rewriting logic (Maude).

**KEYWORDS** Global behavioral consistency, Consistency verification, Multi-modeling, Heterogeneous models, Rewriting Logic, Graph transformation

## 1. Introduction

Model-Driven Engineering (MDE) addresses the increasing complexity of software systems by employing models to describe the different aspects of the system. In this way, MDE promotes a clear separation of concerns and raises the abstraction level throughout the entire development process (France & Rumpe 2007). These models are then used to generate portions of the system leading to an increase in productivity and reduction of errors (Brambilla et al. 2017). As multiple interacting systems are needed to realize the requirements of complex domains, a set of corresponding models would be needed to represent these systems and their interactions. Such a collection of inter-related models is referred to as a *multi-model* (Boronat et al. 2009; Stünkel et al. 2021), which is usually heterogeneous, meaning it consists of models conforming to different modeling languages. Models in a multi-model contradicting each other

can lead to problems during development, system generation, and system execution. Consequently, continuous multi-model consistency management during the development process is a significant issue for multi-models (Spanoudakis & Zisman 2001; Cicchetti et al. 2019).

Recent research describes methods to check the structural consistency of a multi-model (Stünkel et al. 2021; Klare & Gleitze 2019). Structural models, like UML class diagrams, describe structural aspects of systems, i.e., domain concepts and relations between these concepts. This is usually referred to as the denotational semantics of the software system, as it only describes the set of valid instances or states of the system. Nevertheless, approaches to multi-model consistency management must also include a means to maintain *behavioral consistency* since behavioral models, like Business Process Modeling Notation (BPMN) models, are associated with execution semantics describing dynamic aspects of the system (Object Management Group 2017, 2013). Multi-models consisting of different interacting behavioral models are found, for example, when modeling embedded and cyber-physical systems (Vara Larsen et al. 2015).

### JOT reference format:

Tim Kräuter, Harald König, Adrian Rutle, Yngve Lamo, and Patrick Stünkel. *Towards behavioral consistency in multi-modeling*. Journal of Object Technology. Vol. vv, No. nn, 2023. Licensed under Attribution 4.0 International (CC BY 4.0) <http://dx.doi.org/10.5381/jot.2023.vv.nn.a1>

Several approaches exist for checking the consistency of pairs of behavioral models. For example, consistency checking for sequence diagrams and statecharts was implemented using Petri nets (Yao & Shatz 2006) and Communicating Sequential Processes (CSP) (Küster & Stehr 2003). Moreover, some approaches for model simulation in heterogeneous scenarios have been developed, such as Ptolemy (J. Eker et al. 2003) and B-COoL for GEMOC studio (Vara Larsen et al. 2015). However, current approaches either only allow for consistency checking of two behavioral models or do not allow for the definition and checking of global behavioral properties.

We propose a novel approach for consistency management of heterogeneous multi-models, which allows us to define and check *global* behavioral properties. Our approach facilitates specifying *interactions* between multiple potentially heterogeneous behavioral models, which are in turn used to generate a specification of the global behavior. Our approach is generally independent of the particular formalism used in the specification, and currently, we can generate specifications in two different formalisms. The generation of the global behavior specification is *fully automatic* and results in graph transformation rules (or respectively, term rewriting rules) executable in Groove (Maude). Afterward, we can use the built-in verification mechanisms in Groove (Maude) to check the previously defined global behavioral properties.

The remainder of this paper is structured as follows. We introduce a simplified use case (section 2) before explaining our behavioral consistency management approach in detail (section 3). Afterward, we show how we can use the Groove tool set and the Maude system to check behavioral consistency (section 4). Furthermore, we examine related work in section 6. Finally, we discuss the state space explosion problem in section 5 and conclude in section 7.

## 2. Use Case

This section motivates our approach with a simplified use case in which a traffic management system is developed to guide the traffic at a T-Junction with three traffic lights. The traffic management system should control the traffic by switching between the two traffic phases highlighted in figure 1. In addition, it must fulfill the following two requirements. First, it must guarantee safe traffic by changing the three traffic lights, A, B, and C, correctly. Second, it should prioritize arriving buses, i.e., switch the traffic lights quicker than usual to let an approaching bus pass (early green). This so-called bus priority signal is a widely implemented technique to improve service and reduce delays in public transport.

To develop the behavior of the traffic management system, we follow an MDE approach. First, we model the behavior of a traffic light as a Unified Modeling Language (UML) state machine. Then we use BPMN to model the different traffic phases of the T-Junction, including the prioritization of approaching buses.

Using different behavioral modeling languages in the use case has two reasons. First, two software development teams might work on the system in parallel but prefer different model-




Figure 1 Traffic phases of a T-Junction

ing languages. Second, each team can choose the most appropriate modeling language for defining their part of the system. In this use case, the behavior of a traffic light and a T-Junction differs significantly in complexity and requirements, resulting in the use of two different behavioral modeling languages, namely UML state machines, and BPMN.

The behavior of a traffic light is straightforward since it uses only three colors to guide the traffic. Figure 2 shows the typical European<sup>1</sup> traffic light that switches from red to red-amber, green, amber, and back to red. The start state of the traffic light in figure 2 is red but can be any of the four possible states.

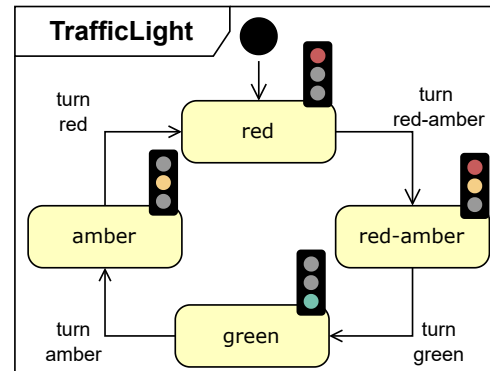


Figure 2 Traffic light state machine model

However, the T-Junction's behavior is more complex since it should coordinate the three traffic lights and communicate with approaching buses to implement bus priority. Consequently, we are using BPMN to model this aspect of the system's behavior and utilize BPMN message and signal events to implement the communication with approaching buses.

We model two processes, one for the T-Junction and one for the Bus. Each process is modeled in its BPMN *pool*. A pool is depicted as a horizontal lane with a name on the left. Message flows (arrows with dashed lines) are only allowed between two different pools.

Figure 3 shows how a possible controller for a T-Junction behaves in the traffic management system. When a TJunction controller is started, we assume that the traffic lights are showing the colors according to phase 1 (see figure 1). Thus,

<sup>1</sup> Traffic lights in other parts of the world might not show red and amber simultaneously before switching to green.

the controller enters a subprocess called phase 1 (see top right in [figure 4](#)), which we describe together with the subprocess called phase 2 later. However, when a fixed amount of time has passed, the subprocess is interrupted by the attached timer boundary event. Then, the controller executes the next activity and switches to phase 2. The controller will pass a throwing signal event before entering a subprocess for phase 2 and repeat the same steps. This signal event represents a broadcast to all buses waiting for traffic light B to become green. After switching back from phase 2 to phase 1 and signaling that traffic lights A and C are green, the controller can stop or execute the described steps again. Typically, the controller does not stop, which is indicated by the default sequence flow going back to the process beginning.

[Figure 4](#) shows the communication of a bus with the subprocess phase 1. The BPMN model and communication for phase 2 of the controller can be defined accordingly.

The phase 1 model uses an event-based gateway to respond to two different kinds of messages. First, the traffic light status can be requested, which is answered by sending a message declaring that the traffic lights A and C are green while B is red. Moreover, early green for traffic light B can be requested. This request ends the subprocess, and the controller immediately switches to phase 2 (see [figure 3](#)), which results in the traffic light B turning green.

The bottom of [Figure 4](#) shows the controller for a bus parameterized with direction B. It will first request the traffic light status to determine if traffic light B is green. If it is green, the bus can pass the junction. However, if it is red, the bus requests to change B to green and waits for a signal that the controller has changed the traffic light. After receiving the signal, the bus passes the junction. A BPMN model for a bus controller parameterized with the direction A or C looks nearly identical. In addition, the bus controller communicates with the phase 2 subprocess, which we only hint at in [figure 4](#). The phase 2 subprocess has the same structure as the phase 1 subprocess but reports that A and C are red, while B is green. Similarly, it terminates if green is requested for A or C. The full model and all other models are available from ([Kräuter 2022](#)).

Having developed behavioral models for the system, we want to check the previously stated *safe traffic* requirement while buses are prioritized. We can lower the overall development cost if we find bugs related to these requirements as early as possible during system development. However, the traffic light model is currently not related to the T-Junction and bus models while the T-Junction is supposed to control the traffic lights, for example, when it switches between the two traffic phases. In addition, the system has to manage multiple slightly different instances of the behavioral models. For example, there are three traffic lights at one T-Junction starting in different states—i.e., showing different colors—and buses approaching the T-Junction from one of the three directions. Consequently, we need a model of the system to allow us to define interactions between the models and configure instances of the behavioral models contained in the multi-model.

The resulting model called the *system relationship model* (SRM) is shown in [figure 5](#) using a graph-based syntax. It con-

tains one node for each behavioral model and arrows to depict behavioral relationships, leading to possible interactions. In addition, it contains enumerations to parameterize the behavioral models. A T-Junction has three associated TrafficLights, A, B, and C, and a set of currently approaching Buses. A TrafficLight has four possible TrafficLightStates and an attribute to define its startState. A Bus has a direction that indicates which TrafficLight of the T-Junction it is approaching.

Finally, using the SRM, we can define a test configuration of our traffic management system to check its requirements. [Figure 6](#) depicts the test system configuration as an instance of the SRM. First, it contains three instances of the traffic light behavioral model, representing the three traffic lights, A, B, and C. Second, it contains an instance of the T-Junction behavioral model connected to the three traffic lights and two instances of the bus behavioral model. Thus, the test system configuration describes a system that controls one T-Junction with three traffic lights and two buses approaching from directions A and B.

First, we would like to check the safe traffic requirement. Since we only want to check system conformance concerning the two traffic phases, we do not need to include the two buses depicted in the red dotted square in [figure 6](#) in the analysis. We cannot simply assert that the system is either in phase 1 or phase 2 since there are intermediate states during the transition between the two phases, which are allowed. By consulting [figure 2](#), we can, for example, expect a state in which traffic lights A and C are amber, and traffic light B is red-amber before reaching phase 2. However, we can define *safe traffic* as the absence of *unsafe traffic*, which is easier to define.

For the T-Junction, unsafe traffic occurs if traffic light A is green or amber and traffic light B is green or amber simultaneously. In addition, the same state combinations are forbidden for traffic lights B and C. Unsafe traffic occurs only in these situations since green and amber mean that cars are allowed to pass, while red (red-amber) means cars are not (not yet, respectively) allowed to pass. We can formalize the consistency requirements as safety properties in Linear Temporal Logic (LTL), i.e., states that should never be reached. The resulting global properties (1) and (2) are the following, assuming the existence of atomic propositions for each traffic light state.

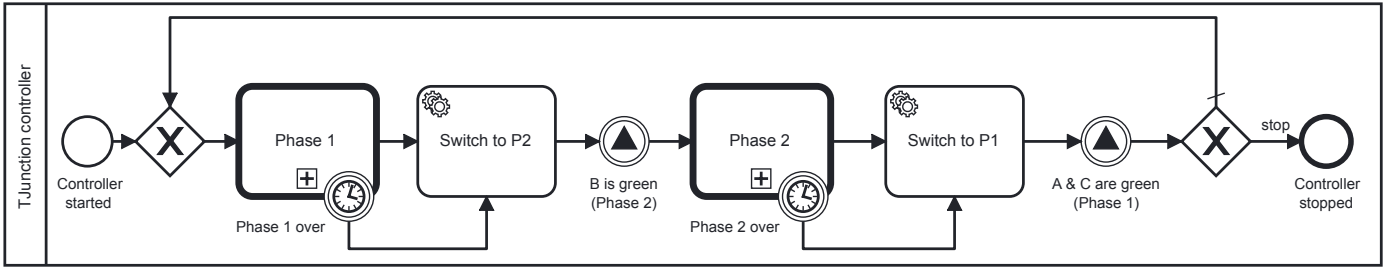
$$\Box \neg ((A_{green} \vee A_{amber}) \wedge (B_{green} \vee B_{amber})) \quad (1)$$

$$\Box \neg ((C_{green} \vee C_{amber}) \wedge (B_{green} \vee B_{amber})) \quad (2)$$

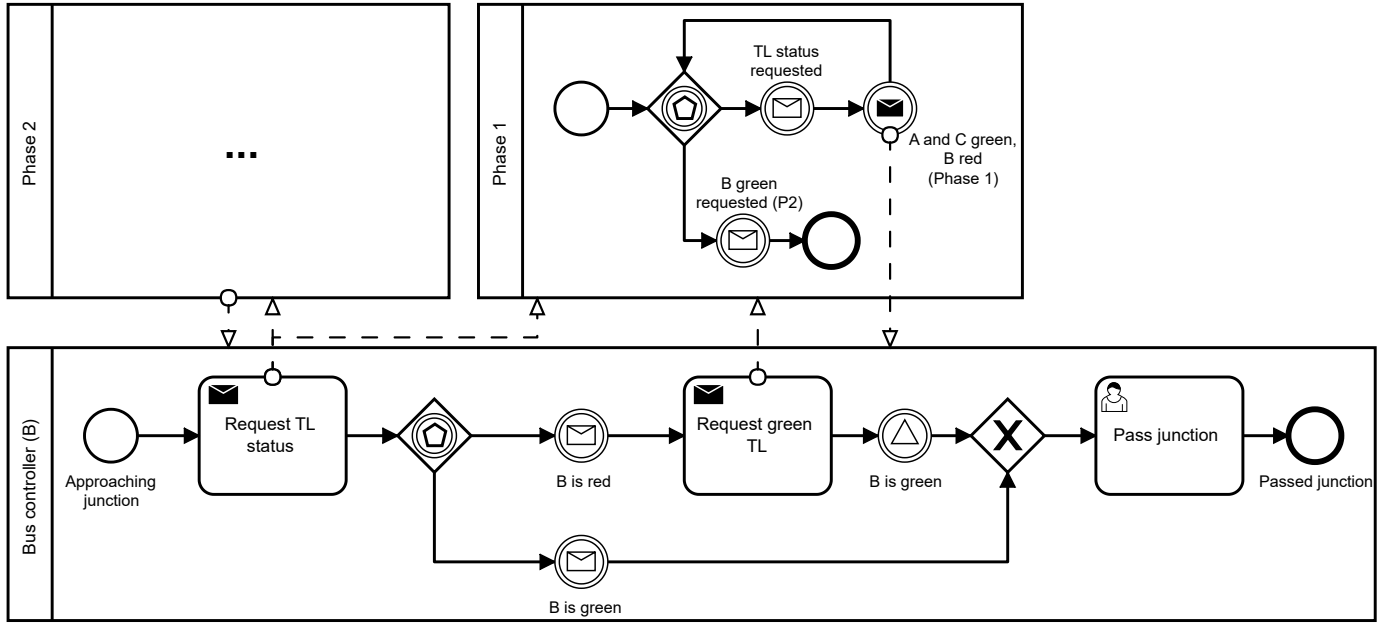
If we include buses B1 and B2 in the system, we would like to check that they cannot pass when their traffic light is red or red-amber. Concretely, this means the *Pass Junction* activity should not execute while the corresponding traffic light is red or red-amber. We formalize these requirements again by using LTL safety properties (3) and (4), where the atomic proposition  $B1_{passing}$  and  $B2_{passing}$  represent that *Pass Junction* (see [figure 4](#)) has started but not finished yet.

$$\Box \neg (B1_{passing} \wedge (A_{red} \vee A_{red-amber})) \quad (3)$$

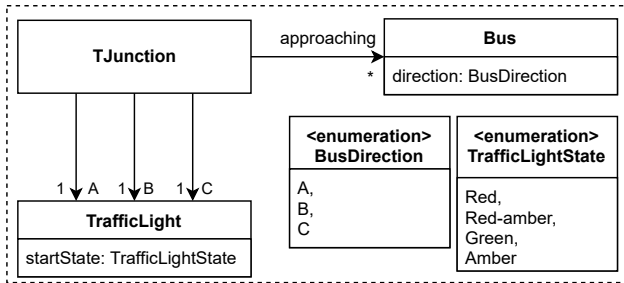
$$\Box \neg (B2_{passing} \wedge (B_{red} \vee B_{red-amber})) \quad (4)$$



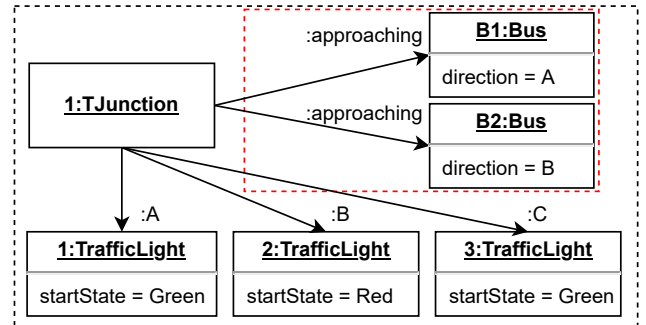
**Figure 3** Model for a TJunction controller



**Figure 4** Model for a bus with direction B and its communication with a T-Junction



**Figure 5** System relationship model of the traffic management system



**Figure 6** Test system configuration

However, to check the global properties, we must execute the system with the behavior specified in the behavioral models according to the test configuration. This is not straightforward since the multi-model of the use case consists of a SRM relating two heterogeneously typed behavioral models. In addition, the system configuration instantiates two behavioral models multiple times with different parameters. Furthermore, we face the problem that the models are not independent of each other.

For example, the T-Junction controller must decide when the traffic lights A, B, and C switch states. Thus, if we were to run the models independently in parallel, the properties would be violated.

A multi-model is behaviorally consistent if it satisfies all of its behavioral properties. A behavioral property is given in temporal logic, for example, LTL in the use case, and is characterized as *local* if it constrains only one model and as *global*



if it spans two or more models in a multi-model. Furthermore, global properties depend on the system configuration, i.e., the instance of the SRM used. In the remainder of this paper, we will describe our approach to address behavioral consistency in multi-modeling and apply it to this use case.

### 3. Behavioral consistency management

Figure 7 depicts our approach to behavioral consistency management as a BPMN diagram.

Our approach consists of five steps.

1. We define a *system relationship model (SRM)* describing which behavioral models may interact.
2. We specify consistency requirements as global behavioral properties for the SRM.
3. We define *interactions* between the behavioral models using the SRM.
4. We automatically generate a specification of the specified global behavior using the interactions and the SRM.
5. Given a system configuration, we check the global behavioral properties using the generated specification.

The behavioral consistency management workflow potentially uncovers violations, describing counter-examples in which the global behavioral properties are not valid. In the following sections, we will describe each step in detail. In addition, the appendix contains figure 18, which gives an overview of all the new concepts and how they are applied to the use case.

#### 3.1. Define the system relationship model

As mentioned, a set of behavioral models might be used to describe the behavior of a software system. Each of these models conforms to its metamodel, corresponding to the behavioral language used to specify the model. The metamodel ensures that models specified in the corresponding languages are well-defined and machine-readable. This is crucial when automating parts of the consistency checking.

The use case utilizes state machine and BPMN models, which conform, respectively, to the metamodels of state machines (see figure 8) and BPMN (see figure 9). The metamodel of state machines is defined by a UML class diagram. In addition, the clouds depict the concrete syntax that we use to denote the models conforming to the metamodel. The traffic light model in figure 2 uses this concrete syntax.

A *StateMachine* has a *startState* and transitions, whereas each *Transition* connects two *States*. The states of a state machine are not explicitly modeled but can be derived from the states connected by the transitions of a state machine, assuming no isolated states.

The metamodel for BPMN (see figure 9) is defined analogously to the one of state machines. A BPMN *Process* contains a set of *FlowNodes* connected by *SequenceFlows*. *FlowNodes* and *SequenceFlows* are *FlowElements*, inheriting an *id* and a *name*. A *FlowNode* can be an *Activity*, *Gateway*, or *Event*. All

special activities, gateways and events are defined in the BPMN specification (Object Management Group 2013).

Behavioral models *interact* to realize the global system behavior. A SRM describes which behavioral models exist in the system and whether they are behaviorally related, i.e., they may interact during execution. In our approach, we define a system relationship metamodel to formally specify these relationships (see figure 10).

We are using a graph-based syntax to define SRMs (see figure 5), where each node corresponds to a behavioral model (typed by a *BehavioralMetamodel*), while each arrow corresponds to a *BehavioralRelationship* (see concrete syntax depicted in clouds). For example, the SRM for the use case (see figure 5) has three behavioral relationships from *TJunction* to *TrafficLight* since a *TJunction* controller interacts with three different traffic lights A, B, and C. Furthermore, there is a behavioral relationship from *TJunction* to *Bus* because we want to check the safety properties (3) and (4). To summarize, behavioral relationships define *which* behavioral models *may* interact, while the interactions in the next step of the workflow describe *how* they interact.

In addition, we allow enumerations and attributes in SRMs. These may be used as parameters, e.g., to define the start state in a state machine (see figure 5). Different instances of the SRM can be used to analyze the global behavior of *different* system configurations by changing the parameters.

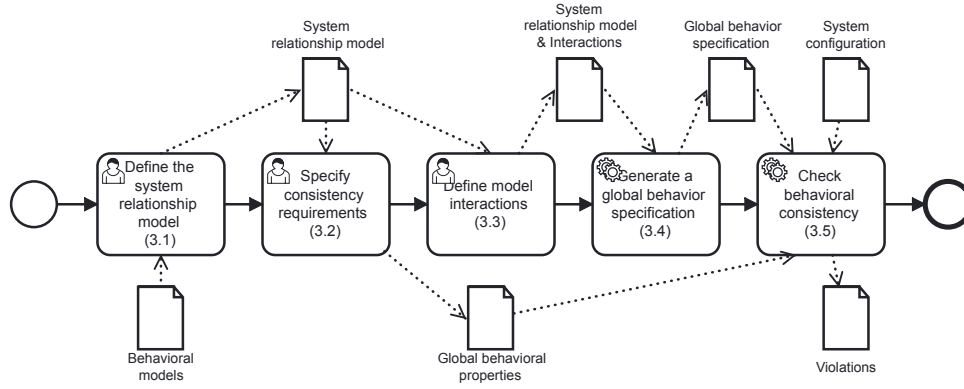
#### 3.2. Specify consistency requirements

In this step, we specify behavioral consistency requirements as global behavioral properties. These properties are defined using a temporal logic, for example, LTL as in the use case. Temporal logic properties are built upon a set of atomic propositions which are either true or false in a given state. Each behavioral language (e.g., the BPMN) specification defines how these states are represented. Furthermore, the transitions between these states depend on the semantics of the behavioral models. In our approach, we call the state representation a *snapshot metamodel* which is specified using class diagrams.

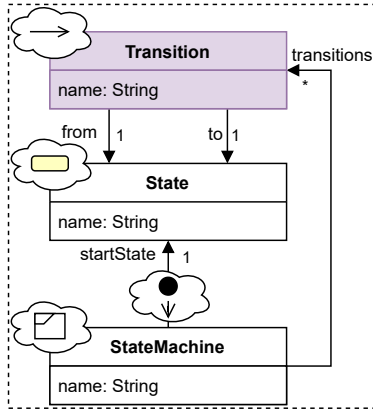
For example, in the use case, we define snapshot metamodels for state machines and BPMN. A state machine is in one state at a time, as shown in the snapshot metamodel on the left of figure 11. We are reusing the concrete syntax elements from the state machine metamodel (see figure 8) for the snapshot metamodel. In addition, each snapshot metamodel has a root element in our approach, highlighted in light blue.

The snapshot metamodel for BPMN is based on a *Token distribution* as described in the BPMN specifications (Object Management Group 2013) (see on the right of figure 11). The root element *ProcessSnapshot* has tokens and subprocesses. A *Token* indicates where it is located in the BPMN model using its *position* attribute. A valid position is the id of a *FlowElement* (see figure 9). Also, for the snapshot metamodel of BPMN we reuse the concrete syntax of the BPMN metamodel. In addition, *Tokens* are highlighted with green bubbles in the middle of sequence flows and the top right of an activity.

Using the snapshot metamodels, one can only specify local behavioral properties. To specify global behavioral properties,



**Figure 7** Behavioral consistency management workflow



**Figure 8** Finite state machine metamodel


we combine the information of the SRM with the snapshot metamodels. Each behavioral model is typed by a behavioral metamodel, which has a snapshot metamodel describing its state structure. Thus, we know how states of behavioral models are represented when they are instantiated. For example, [figure 12](#) shows how to specify the atomic propositions  $A_{green}$  and  $B1_{passing}$ , used in the global properties in [section 2](#). Snapshot links connect instances of behavioral models with root element instances of the corresponding snapshot metamodel.

To make formulating atomic propositions less cumbersome, one can use the concrete syntax of the individual snapshot metamodels. For example, [figure 13](#) shows the same atomic propositions as [figure 12](#) but uses the introduced concrete syntax.


With the defined atomic propositions as ingredients, one can use temporal logic to define global behavioral properties such as the properties (1)-(4) in [section 2](#). It is worth noting that the defined atomic propositions are *model-specific*, meaning they exactly fit the given multi-model.

### 3.3. Define model interactions

Structural models in a multi-model might contain related information. Thus, current approaches define so-called *commonalities* to explicate these relationships and keep the information consistent ([Stünkel et al. 2021](#); [Klare & Gleitze 2019](#)).



**Figure 9** Simplified BPMN metamodel ([Object Management Group 2013](#))



**Figure 10** System relationship metamodel excerpt

To check the behavioral consistency of a system, we are primarily interested in global system behavior. However, global behavior depends not only on the local behavior of the models but also on their interactions. Consequently, we call these inter-model relationships *interactions* since they carry behavioral meaning while commonalities carry structural meaning ([Kräuter 2021](#)).

To specify interactions between different behavioral models, we defined an interaction language given by the full system relationship metamodel in [figure 14](#). In addition, we introduce the concept of *state-changing elements*. Each BehavioralMetamodel specifies a set of StateChangingElements. For example, a state machine defines states and transitions, but only the transitions describe how the states in a state machine change. Thus, the transitions are the StateChangingElements of a state machine (highlighted in purple, see [figure 8](#)). Similarly, the

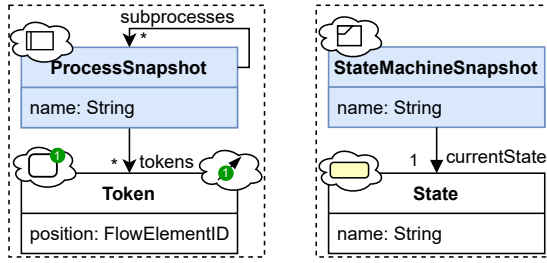


Figure 11 BPMN and FSM snapshot metamodels

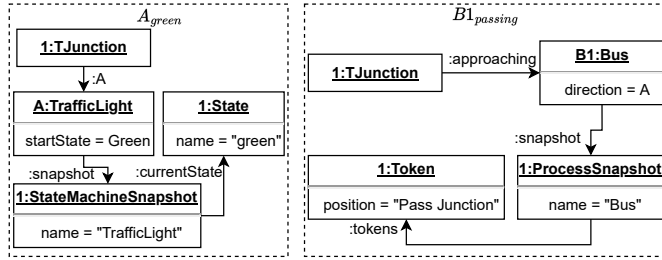


Figure 12 Atomic propositions  $A_{green}$  and  $B1_{passing}$

flow nodes are the StateChangingElements of a BPMN process (highlighted in purple, see figure 9)<sup>2</sup>. The interaction of behavioral models is only possible through instances of the StateChangingElements.

Our approach is based on the requirement that state-changing elements can be identified in metamodels for any behavioral formalism. This requirement is not difficult to meet since behavioral modeling languages must have some *observable* construct to describe state changes. Inspecting other behavioral languages, such as Petri Nets or UML activity diagrams, shows that identifying state-changing elements (transitions, activity nodes) is unproblematic.

An Interaction has a sender, a set of receivers, and an InteractionType. Currently, there is only the synchronous InteractionType. However, more interaction types, for example, asynchronous interactions or interactions with message passing, could be added in the future. We model the role of sender and receiver in interactions to accommodate these interaction types in the future. The two roles are not needed for synchronous interactions and are hidden from the user in the concrete syntax used later in listing 1. Furthermore, asynchronous interactions can be modeled using two synchronous interactions with an additional behavioral model, such as a queue. Each InteractionReceiver references one BehavioralRelationship and one StateChangingElement.

The sender and the elements of the InteractionReceivers describe a state change for their behavioral models. By connecting them with a synchronous interaction, we define simultaneous state changes in one atomic step. Consequently, an interaction can define a synchronization between behavioral models. In addition, one can only define interactions for state-changing elements if their behavioral models are connected by a behavioral

<sup>2</sup> One exception is the event-based gateway, which is not part of the StateChangingElements.

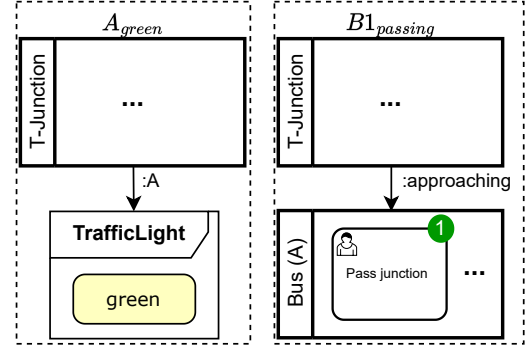


Figure 13 Concrete syntax for  $A_{green}$  and  $B1_{passing}$

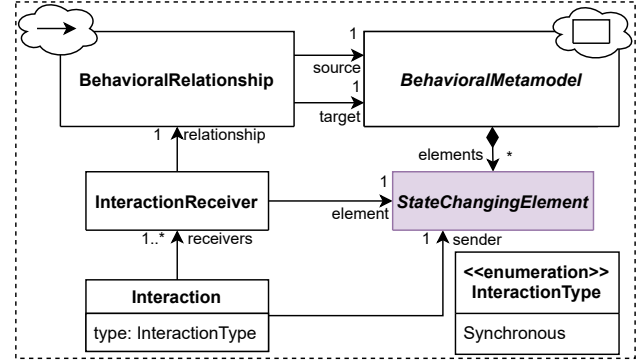


Figure 14 System relationship metamodel

relationship (see constraint (5)). We use "." in constraints to navigate along associations. For example, i.receivers means following all receivers links for an Interaction object, resulting in a set of InteractionReceiver objects.

$$\begin{aligned} \forall i \in \text{Interaction} : \forall r \in i.\text{receivers} : \\ r.\text{element} \in r.\text{relationship}.\text{target}.\text{elements} \wedge \\ i.\text{sender} \in r.\text{relationship}.\text{source}.\text{elements} \end{aligned} \quad (5)$$

We allow the definition of as many interactions as desired. The interactions restrict the local behavior of the models since some parts of their behavior must be synchronized. Two interactions are not allowed to share state-changing elements (see constraint (6)). If such a situation occurs, it must be resolved by the modeler by deleting one of the interactions or merging the two interactions into one.

$$\begin{aligned} \forall i_1, i_2 \in \text{Interaction} : \\ (i_1.\text{receivers}.\text{element} \cup i_1.\text{sender}) \cap \\ (i_2.\text{receivers}.\text{element} \cup i_2.\text{sender}) = \emptyset \end{aligned} \quad (6)$$

In the use case, the TJunction controller interacts with the three traffic lights, A, B, and C. Listing 1 defines two interactions synchronizing TJunction controller and the traffic lights, using a textual Domain-Specific Language (DSL). The **synchronize** keyword specifies the InteractionType to be synchronous. Each interaction first defines the *sender* state-changing element. Then *receivers* are defined by navigating along an arrow (BehavioralRelationship) to other BehavioralModels in the

SRM and then specifying a StateChangingElement. All navigation along BehavioralRelationships starts from the BehavioralModel containing the *sender*, so constraint (5) is fulfilled.

```

1 synchronize(TJunction.Switch_to_P1 ,
2             TJunction.A.turn_green ,
3             TJunction.B.turn_red ,
4             TJunction.C.turn_green)
5
6 synchronize(TJunction.Switch_to_P2 ,
7             TJunction.A.turn_red ,
8             TJunction.B.turn_green ,
9             TJunction.C.turn_red)

```

**Listing 1** Interactions for the use case

We can explain the interactions as follows. The first interaction defines that the task Switch to P1 and three other state-changing elements synchronize. Furthermore, line 2 specifies that one of the synchronization receivers is the element *turn green* connected by the relationship A. Similarly, two other transitions of the traffic lights B and C are specified in the following two lines. Thus, the interaction defines a synchronization of a task and three traffic light transitions. The second interaction defines a synchronization for the task Switch to P2 and three traffic light transitions.

To summarize, we use the system relationship metamodel to define the relations between the behavioral models in a multi-model. Thus, the inter-relations between behavioral models in a multi-model are given by interactions and their behavioral relationships.

### 3.4. Generate a global behavior specification

Using the SRM and snapshot metamodels, we can represent the global states of the system. However, we still need a formal specification of the global behavior to check the defined properties. The specification of the global behavior used in our approach must fulfill the following three requirements:

1. The specification must respect the semantics of each behavioral model, except for interactions.
2. The specification must reflect the defined interactions between the behavioral models.
3. The specification semantics must allow the checking of behavioral properties for a given system configuration.

Thus, a specification in any formalism fulfilling these three requirements can be used in our approach. Consequently, one can experiment with different formalisms, for example, graph transformations, rewriting logic, state machines, Petri nets, or process algebras, without changing the general framework. One can then pick the most suitable formalism for the modeling scenario at hand regarding, for example, the performance of consistency checking. In section 4, we describe how we generate specifications for the graph transformation toolset Groove and the term-rewriting system Maude.

To summarize, we generate a specification of the global system behavior. This generation takes the models and interactions as input and is fully automated to allow frequent model changes.

### 3.5. Check behavioral consistency

In this step, we use the generated specification of the global behavior to check consistency given a system configuration. A system configuration is an instance of the SRM and is automatically translated into the formalism used in the specification. We then check the consistency of the defined properties using the specification and the system configuration. This step is fully automated, such that it can be executed as many times as needed for different system configurations and properties but using the same specification.

Finally, if a consistency requirement is not fulfilled, a violation, including a counterexample, will be presented to the user. Adopting the same concrete syntax to visualize the counterexample as for the atomic propositions should be ideal for helping a user understand it. An unsuccessful consistency check leads to a consistency restoration process, which is crucial but out of the scope of this paper. We describe consistency checking for the use case and its result at the end of the next section.

## 4. Specification of the global behavior

In this section, we will describe our implementation to generate formal specifications in Groove (Rensink 2004) and Maude (Clavel et al. 2007). Then, we use the generated specifications to check behavioral consistency for the use case and discuss the results.

Both implementations utilize a *global higher-order model transformation (MT)* from the behavioral models and their interactions to graph transformation (GT) rules (Groove) or rewriting rules (Maude). Since the results of the MT can be regarded as MTs themselves, we say the MT is *higher-order* (Tisi et al. 2009). We will now describe this transformation using the term *rules* to mean either GT or rewriting rules. The details specific to each implementation are given in the following sections.

The higher-order MT can be decomposed into two steps. The first step to generate a specification of the global behavior is to generate rules for each behavioral model contained in the multi-model. Each set of rules must describe the behavior of the given behavioral model by manipulating instances of the snapshot metamodels. For example, a rule for a transition in a state machine changes the current state of a state machine snapshot from the source to the target of the transition.

Thus, we need *local higher-order MTs* for each behavioral modeling language producing sets of rules. Each local higher-order MT only has to be implemented once by an MDE tool developer and made available, for example, as a plugin, such that it can then be reused in any future setting the language is needed. In addition, each local higher-order MT must keep traces of the generated rules. Concretely, it has to save which rules originated from which state-changing elements in the behavioral model. Returning to the state machine example, we must know which transition results in which rule. In general, multiple rules may be associated with one state-changing element of a behavioral model. For example, a receive task in a BPMN process is represented by two rules since it starts and then waits for an incoming message before finishing.

The second step is to modify the generated rules to reflect



the defined interactions. Interactions define the synchronization of systems which we encode by merging the individual rules into rules describing the global behavior. The merging process to obtain the global rules is done in two steps:

1. Rules generated from state-changing elements that are *not* part of interactions remain unchanged and are added to the global rule set.
2. For each interaction, we do the following:
  - (a) Find the corresponding rule<sup>3</sup>  $P_0$  for the sender of the interaction and find the rules  $P_1, P_2, \dots, P_n$  for the receiver state-changing elements using the saved traces.
  - (b) Create a *global rule* for the rules  $P_0, P_1, \dots, P_n$ , which applies all of them at once, i.e., synchronizes the state changes of the behavioral models.
  - (c) For each receiver of the interaction, instantiate the corresponding behavioral relationship from the behavioral model in  $P_0$  to the behavioral model in  $P_i$ , for  $1 \leq i \leq n$ , and add it to the global rule. Thus, only behaviorally related models may interact, i.e., change their state simultaneously.

Figure 17 shows the global rule resulting from the merging process for graph transformation rules. The global rule contains four individual rules changing the state of three traffic lights state machines and one T-Junction BPMN process due to step 2(b). In addition, it contains three instances of behavioral relationships connecting the behavioral models as described in step 2(c). In the following sections, we describe in more detail how global rules for interactions are created in Groove and Maude.

#### 4.1. Groove specification

This section describes how graph transformations can be used as one possible formalism for behavioral consistency management. We utilize the Groove tool set to run the generated specifications, i.e., GT systems (Rensink 2004).

A GT system consists of a set of GT rules of the form  $L \rightarrow R$ , where the graph  $L$  is called the left-hand side and the graph  $R$  is called the right-hand side of the rule. Nodes/edges in  $R$  but not in  $L$  are added by a rule, while nodes/edges in  $L$  and  $R$  are preserved, and a rule deletes nodes/edges that are in  $L$  but not in  $R$ . Applying a GT system to a given graph, one obtains a state space where each state is a graph, and each transition is a rule application. A formal description of GT systems can, for example, be found in (Ehrig et al. 2006).

We generate typed GT systems, where the merge of the SRM and the snapshot metamodels is the type graph (Kräuter 2022). Individual rules and rules changing the global state conform to this type graph. Interactions result in global rules which change multiple parts of the global state. A global GT rule is calculated by taking the sum of all left-hand sides and right-hand sides of the individual GT rules (implements step 2(b)

<sup>3</sup> If one state-changing element results in more than one rule, one can define a strategy to pick the appropriate rule.

in section 4 for Groove) (Baldan et al. 1999, Definition 3.2.7). Together with a system configuration, i.e., a start graph, we can obtain an executable formal specification of the global behavior. Consequently, this can be used to check behavioral consistency.

We will now explain how the GT rules are generated for the use case. To apply our approach to the use case, we need to define local higher-order MTs from state machines and BPMN processes to GT rules.

**4.1.1. State machine semantics** The local higher-order MT to generate GT rules for finite state machines is straightforward. Each transition leads to a GT rule. For example, figure 15 shows the GT rules for the transition turn red-amber of the traffic light model. It uses the concrete syntax introduced in figure 11 to depict state machine snapshots and their current states. We depict a GT rule by showing the graph  $L$  on the left,  $R$  on the right, and a named white arrow from  $L$  to  $R$ .




Figure 15 GT rule for *turn red-amber*

Using a traffic light snapshot with the state red as a start graph, we generate the same state space in Groove as the traffic light state machine describes. All generated GT systems, including further instructions regarding execution and consistency checking, can be found in (Kräuter 2022).

**4.1.2. BPMN semantics** The local higher-order MT to generate GT rules for BPMN processes is challenging, and we are currently only supporting a subset of the BPMN semantics (Kräuter et al. 2022). Generally, we construct one or more rules for each flow node in a BPMN model. Figure 16 shows the GT rules for the task Switch to P1 of the TJunction controller. It uses the concrete syntax introduced in figure 11 to represent process snapshots containing tokens.


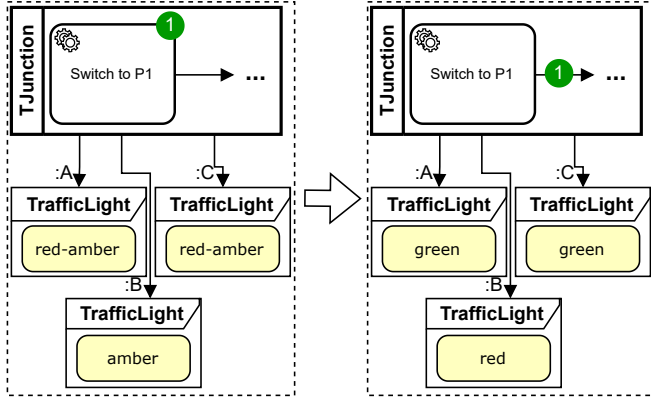


Figure 16 Example GT rules for the TJunction controller

Due to limited space, we only show the Switch to P1 rules but the artifacts of this paper (Kräuter 2022) contain the full GT system. A detailed explanation of the GT rule generation is given in (Kräuter et al. 2022). To summarize, we can generate GT systems that implement the behavioral semantics of BPMN.

**4.1.3. Check behavioral consistency** The defined interactions change the rules for switching to phases 1 and 2. Figure 17 shows the resulting rule for switching to phase 1. We decided that the interactions between the traffic lights and the TJunction controller should synchronize with the end of the task, not the start.



**Figure 17** Global GT rule to switch to phase 1

The rule changes all traffic lights simultaneously and finishes the task. The corresponding individual rules no longer exist, so synchronization of the behavior is guaranteed. A similar rule exists for switching to phase 2, resulting from the other interaction. All other rules are left untouched while constructing the global GT system.

Finally, we can generate the global state space of the system using the global GT system, which can be found in (Kräuter 2022). To check the requirements formalized by properties 1-4, we must also encode the used atomic propositions. These are specified as *graph conditions* in Groove. A graph condition in Groove is a rule that does not change elements but can be used as an atomic proposition in model checking.

## 4.2. Maude specification

We are using the predefined Maude module for object-based programming to encode the snapshot metamodels in Maude. Thus, we can represent state machine snapshots and BPMN process snapshots as objects in Maude. In addition, we add behavioral relationships from the SRM such that we can instantiate them to create the system configuration described in the use case.

Our Maude rules always rewrite multisets of objects, where individual objects are denoted in angle brackets ( $\langle \dots \rangle$ ). Interactions lead to rule merges in Maude, where the left and right sides of each rule are added to a global rule (implements step 2(b) in section 4 for Maude). The global rule is a valid Maude rule since the union of two multisets always results in a multiset. Furthermore, behavioral relationships, as defined in the interaction, are added to the rule as additional objects on both sides. Similarly to the process for GT rules, individual rules are deleted.

We will now explain how rewriting rules are generated for the use case. To apply our approach to the use case, we need to define local higher-order MTs from state machines and BPMN processes to rewriting rules.

**4.2.1. State machine semantics** The generation of rewriting rules is similar to the generation of GT rules. Each state machine transition leads to one rewriting rule, which rewrites state machines conforming to the state machine snapshot metamodel. Listing 2 shows the rewriting rule for the transition turn red-amber of the traffic light model. The rewriting rule is equivalent to the GT rule in figure 15.

Using Maude's search command and search graph visualization, we can verify that the generated state space conforms to the traffic light state machine. The generated Maude specifications and example executions can be found in (Kräuter 2022).

**4.2.2. BPMN semantics** The rewriting rule generation is similar to the GT rule generation and supports the same subset of BPMN semantics. Similarly, one or more rules are constructed for each flow node in a BPMN process, where objects conform to the BPMN snapshot metamodel. Listing 3 shows the Maude rewriting rules for the task Switch to P1 of the TJunction controller. The rewriting rules are equivalent to the GT rules in figure 16.

Like the GT rules, the rewriting rules move tokens by changing their position according to the given BPMN model. Thus, we can generate Maude rewriting rules that implement the behavioral semantics of the BPMN models in the use case (see (Kräuter 2022)).

**4.2.3. Check behavioral consistency** The defined interactions change the rules for switching to phases 1 and 2. Listing 4 shows the resulting rule for switching to phase 1, where synchronization with the end of the task was chosen. The rewriting rule is equivalent to the GT rule in figure 17. The rule changes all traffic lights simultaneously and finishes the task. Similarly, a global rule for switching to phase 2 is created and all individual rules are deleted.

Finally, we can use the Maude model checker module to check the requirements formalized by the LTL properties 1-4. Atomic propositions can be represented in Maude using the encoded SRM and the snapshot metamodels.

## 4.3. Behavioral consistency in the use case

Both specifications show that properties (1) and (2) hold, while properties (3) and (4) do *not* hold (see artifacts in (Kräuter 2022)). The counterexamples for properties 3 and 4 show an unexpected race condition that must be handled: After the TJunction controller signals that the traffic light A/B is green, bus B1/B2 can advance to the Pass Junction activity. However, simultaneously, the TJunction controller can enter the subprocess for the next phase, which can be interrupted by the associated timer event. This can happen before the bus B1/B2 passes the junction, resulting in an invalid state.

The system developers have different options to handle the detected inconsistencies. One option is to keep the models unchanged and pay special attention to the found race condition during system implementation. This can be an acceptable solution since the Pass Junction activity is also modeled as a user activity, i.e., the bus driver decides when to cross the T-Junction. Furthermore, tolerating inconsistencies can be a viable option in MDE (Weidmann et al. 2021). Another option is

```

1 var X : String . --- Object id
2 rl [turn_red_amber] : < X : FSM | name : "trafficLight", state : "red" >
3                     => < X : FSM | name : "trafficLight", state : "red-amber" > .

```

**Listing 2** Maude rule for *turn red-amber*

```

1 vars o0 : Did . --- Object ids
2 vars SIG M T : MSet . --- Signals, messages and tokens
3 vars P S : Configuration . --- Processes and subprocesses
4 --- Switch to P1 start rule
5 rl [Switch_to_P1_start] :
6 < "use-case-execution" : BPMNSystem | messages : (M), processes : (
7   < o0 : ProcessSnapshot | name : "T-Junction controller", tokens : ("Phase_2_Switch_to_P1" T),
8     signals : (SIG), subprocesses : (S), state : Running > P) >
9   =>
10 < "use-case-execution" : BPMNSystem | messages : (none), processes : (
11   < o0 : ProcessSnapshot | name : "T-Junction controller", tokens : ("Switch_to_P1" T), signals : (
12     none), subprocesses : (S), state : Running > P) > .
13 --- Switch to P1 end rule
14 rl [Switch_to_P1_end] :
15 < "use-case-execution" : BPMNSystem | messages : (M), processes : (
16   < o0 : ProcessSnapshot | name : "T-Junction controller", tokens : ("Switch_to_P1" T), signals : (
17     SIG), subprocesses : (S), state : Running > P) >
18   =>
19 < "use-case-execution" : BPMNSystem | messages : (none), processes : (
20   < o0 : ProcessSnapshot | name : "T-Junction controller", tokens : ("Switch_to_P1_A_&
21     _C_are_green_Phase_1" T), signals : (none), subprocesses : (S), state : Running > P) > .

```

**Listing 3** Example rewriting rules for the TJunction controller

to change the models to resolve the inconsistency. For example, the T-Junction controller could wait for the bus to pass before changing the traffic lights again.

## 5. State space explosion

State space explosion is one of the predominant issues when applying model checking to complex systems, which are often found in real-world applications. As one can see in [table 1](#) and [table 2](#), our approach is not immune to state space explosion. The tables show the states, transitions (rewrites), and average runtime of a full state space exploration in the Groove (Maude) specification. Four scenarios were benchmarked, starting with the use case multi-model without approaching buses and then increasing the number of buses up to three.

To calculate the average runtime, we used the hyperfine benchmarking tool ([Peter 2022](#)) (version 1.15.0), which ran state space exploration for each scenario ten times. Timing evaluations were done on a notebook with an AMD Ryzen PRO 3500U processor and 16 GB of RAM running Maude version 3.2.1 (inside the Windows Subsystem for Linux) and Groove version 5.8.1. A description of how to run the benchmarks is available in ([Kräuter 2022](#)).

The benchmarks include start-up times, for example, the state space exploration for the No Buses scenario takes ~500ms according to Groove. Thus, most of the exploration time in Groove and Maude for the No Buses scenario is attributed to the start-up time. Generally, the Maude exploration time is lower despite larger state spaces due to technical differences in implementing the BPMN semantics. However, as expected, neither of the tools is immune to state space explosion.

As in our use case, one is generally not interested in a full state space exploration as in [table 1](#) and [table 2](#) but rather in the validity of a set of behavioral properties. Checking a prop-

Use case	States	Transitions	Exploration time
No buses	168	438	~2.948 ms
1 bus	2.888	10.046	~4.510 ms
2 buses	27.880	121.554	~11.629 ms
3 buses	195.336	1.028.340	~63.487 ms

**Table 1** State space exploration in Groove

Use case	States	Rewrites	Exploration time
No buses	168	664	~63.3 ms
1 bus	3.304	19.958	~277.8 ms
2 buses	35.280	27.9776	~2.994 ms
3 buses	260.176	2.522.582	~28.091 ms

**Table 2** State space exploration in Maude

erty specified in LTL does not necessarily lead to a full state space exploration. Furthermore, not every property is concerned with all behavioral models. For example, properties (1) and (2) of the use case do not involve buses and can be checked on smaller state spaces, not including approaching buses. In addition, checking a set of properties can be run in parallel. If some properties are computationally expensive, they can be run on dedicated hardware, for example, during a continuous integration pipeline once a day in case of a behavioral model or interaction change. Thus, checking the consistency of a multi-model can be seen as an additional test during MDE, which can be run locally but, in addition, is a vital part of continuous

```

1 vars X Tl1 Tl2 Tl3 : String . --- Object ids
2 vars P S : Configuration . --- Processes and subprocesses
3 vars SIG M T : MSet . --- Signals, messages and tokens
4 rl [Switch_to_P1] : < "use-case-execution" : BPMNSystem | messages : (M), processes : (< X :
   ProcessSnapshot | name : "T-Junction controller", tokens : ("Switch_to_P1" T), signals : (SIG),
   subprocesses : (S), state : Running > P) >
5     < "A" : BehavioralRelationship | from : X, to : Tl1 >
6     < "B" : BehavioralRelationship | from : X, to : Tl2 >
7     < "C" : BehavioralRelationship | from : X, to : Tl3 >
8     < Tl1 : FSM | name : "trafficLight", state : "red-amber" >
9     < Tl2 : FSM | name : "trafficLight", state : "amber" >
10    < Tl3 : FSM | name : "trafficLight", state : "red-amber" >
11    =>
12    < "use-case-execution" : BPMNSystem | messages : (none), processes : (< X :
   ProcessSnapshot | name : "T-Junction controller", tokens : ("Switch_to_P1_A_&_C_are_green_Phase_1"
   T), signals : (none), subprocesses : (S), state : Running > P) >
13    < "A" : BehavioralRelationship | from : X, to : Tl1 >
14    < "B" : BehavioralRelationship | from : X, to : Tl2 >
15    < "C" : BehavioralRelationship | from : X, to : Tl3 >
16    < Tl1 : FSM | name : "trafficLight", state : "green" >
17    < Tl2 : FSM | name : "trafficLight", state : "red" >
18    < Tl3 : FSM | name : "trafficLight", state : "green" > .

```

**Listing 4** Global Maude rule to switch to phase 1

integration.

The performance of the Maude LTL model checker is comparable to the popular SPIN model checker (S. Eker et al. 2004) and thus is proven to be competitive. However, one could use different techniques to mitigate the state space explosion problem further. One technique is to abstract models further such that they only contain information relevant to the set of properties to be checked. Thus, minimal models are synchronized, leading to smaller state spaces. However, finding correct minimal models might not be trivial.

Partial-order reduction is a well-known and effective technique to mitigate the state space explosion problem. It is currently not implemented in the Maude LTL model-checker, but there is promising work to integrate partial-order reduction into the model-checker (Farzan & Meseguer 2007), showing substantial state space reductions. The potential to reduce the state space using this technique is huge, especially for model-checking concurrent systems (Clarke et al. 2018). Our approach analyzes concurrent systems with some interaction, and thus model-checking would greatly benefit from partial-order reduction. In our opinion, partial-order reduction must be implemented to analyze models from real-world applications.

## 6. Related work

This contribution builds on our previous publication (Kräuter 2021). However, we refined the behavioral consistency management workflow into five steps and introduced new concepts such as interactions, behavioral relationships, and state-changing elements (see SRM). Moreover, we outlined a general approach for defining atomic propositions, leading to global behavioral properties. In addition, our previous work did not support BPMN models and was solely based on GTs.

The general idea of transforming different behavioral formalisms to a single formalism to reason about cross-cutting concerns is not new, see, e.g., (Engels et al. 2001). For example, (Küster & Stehr 2003) developed consistency checking for se-

quence diagrams and statecharts based on CSP, while (Yao & Shatz 2006) used Petri nets for the same scenario. Furthermore, (Cunha et al. 2011) analyzes sequence diagrams in the context of embedded systems using a transformation to Petri nets. Nevertheless, all approaches focus on a single modeling language or a fixed combination of two languages. Consequently, they do not consider the general problem of behavioral consistency in a heterogeneous modeling scenario.

(Kienzle et al. 2019) proposes a unifying framework for homogeneous model composition. To combine behavioral models, they use Event structures as an underlying formalism and show how different homogeneous behavioral models can be combined. Nonetheless, they do not apply their approach to heterogeneous models. Generally, their approach is compatible with ours since we do not mandate a specific formalism in our approach. However, we have chosen GT rules and term rewriting rules as formalisms to use in our approach since they operate on a higher abstraction level. In addition, it was shown that many formalisms, for example, the  $\pi$ -calculus (Gadducci 2007), can be implemented using GT rules.

(Vara Larsen et al. 2015) tackles the problem of coordinating heterogeneous behavioral models. They coordinate the models using their coordination language B-COoL, similar to the interactions we define between behavioral models. To execute the models with the specified coordination, they transform them into Clock Constraint Specification Language (CCSL) models. Their work also results in plugins for GEMOC studio, which support running and debugging the models. (J. Eker et al. 2003) propose an actor-oriented solution to the model composition problem in the presence of heterogeneity. Their approach is implemented in the tool Ptolemy.

However, both approaches focus on simulation rather than consistency or model-checking. Neither provides concrete means to define atomic propositions and check global behavioral properties. Furthermore, (Vara Larsen et al. 2015) assumes a fixed system configuration while we can check different system configurations.



## 7. Conclusion and future work

Our work represents a first formalization of behavioral consistency management in a heterogeneous modeling scenario, enabling formulating and checking *global* properties. Previous work either only dealt with the behavioral consistency between specific pairs of models (Yao & Shatz 2006; Küster & Stehr 2003) or focused on the simulation in a heterogeneous scenario (J. Eker et al. 2003; Vara Larsen et al. 2015).

Our approach follows five key steps (see figure 7). First, we define which behavioral models in the multi-model *may interact* by creating a SRM. Second, we specify consistency requirements as *global behavioral properties*. Third, we specify how behavioral models interact. Then, we generate a formal specification of the *global behavior* using the SRM and defined interactions. The global behavior specification is based on the defined interactions but preserves the original behavior of each model except for interactions. Finally, we check the *global behavioral properties* using this specification. In addition, not only fully automatized model-checking but also a user-driven simulation of the multi-model and its interactions is possible.

Our approach is independent of the particular formalism used in the generated specification, and we currently support generating GT systems (Groove) and rewriting logic (Maude). The successful implementation shows that our approach can work with different underlying formalisms and serves as a *proof of concept*. Furthermore, we discussed the state space explosion problem and how it can be mitigated.

In future work, we plan to extend our implementation to support more behavioral formalisms such as activity diagrams, hierarchical state machines, and the  $\pi$ -calculus. Especially integrating the  $\pi$ -calculus, which was formalized using GT systems in (Gadducci 2007), will be beneficial since it is profoundly different from the currently supported formalisms.

Furthermore, two systems often exchange data while interacting, for example, using name-passing or messaging. The exchanged data then greatly influences the future behavior of the systems. Thus, adding data transfers to interactions between heterogeneous models is an important issue left for future work.

In addition, it is important to check that running systems behave as specified in their behavioral models, including the defined interactions between them. Thus, extending runtime verification techniques to include systems specified by a multi-model with interactions constitutes a challenging problem.

Finally, if consistency violations are found, consistency restoration must be achieved. We leave consistency restoration of behavioral models as a problem for future work.

## References

- Baldan, P., Corradini, A., Montanari, U., Rossi, F., Ehrig, H., & Löwe, M. (1999, August). Concurrent semantics of algebraic graph transformations. In *Handbook of Graph Grammars and Computing by Graph Transformation* (Vol. 3, pp. 107–188). World Scientific. doi: 10.1142/9789812814951\_0003
- Boronat, A., Knapp, A., Meseguer, J., & Wirsing, M. (2009). What Is a Multi-modeling Language? In A. Corradini & U. Montanari (Eds.), *Recent Trends in Algebraic Develop-*

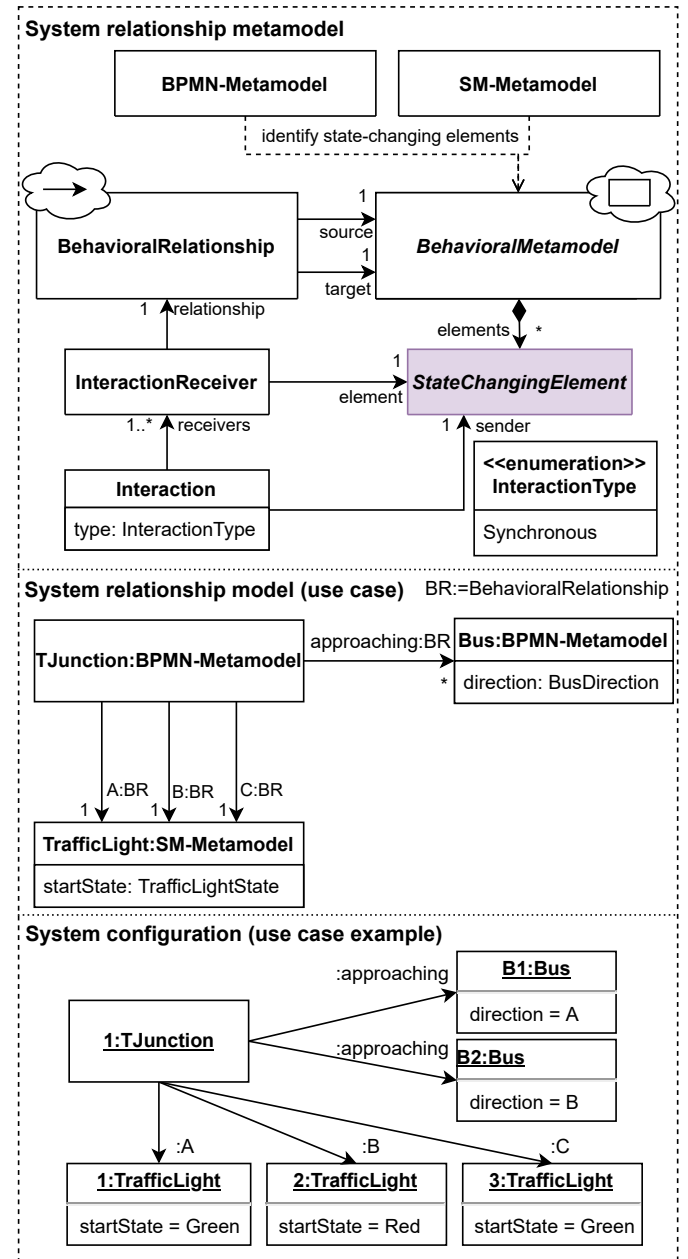


Figure 18 Overview of all concepts and their example usage

ment Techniques (Vol. 5486, pp. 71–87). Berlin, Heidelberg: Springer Berlin Heidelberg. doi: 10.1007/978-3-642-03429-9\_6

- Brambilla, M., Cabot, J., & Wimmer, M. (2017). *Model-driven software engineering in practice* (Second ed.) (No. 4). San Rafael, Calif.: Morgan & Claypool Publishers.
- Cicchetti, A., Ciccozzi, F., & Pierantonio, A. (2019, December). Multi-view approaches for software and system modelling: A systematic literature review. *Software and Systems Modeling*, 18(6), 3207–3233. doi: 10.1007/s10270-018-00713-w
- Clarke, E. M., Henzinger, T. A., Veith, H., & Bloem, R. (Eds.). (2018). *Handbook of Model Checking*. Cham: Springer International Publishing. doi: 10.1007/978-3-319-10575-8
- Clavel, M., Durán, F., Eker, S., Lincoln, P., Martí-Oliet, N.,

- Meseguer, J., & Talcott, C. (2007). *All About Maude - A High-Performance Logical Framework* (Vol. 4350). Berlin, Heidelberg: Springer Berlin Heidelberg. doi: 10.1007/978-3-540-71999-1
- Cunha, E., Custodio, M., Rocha, H., & Barreto, R. (2011, November). Formal Verification of UML Sequence Diagrams in the Embedded Systems Context. In *2011 Brazilian Symposium on Computing System Engineering* (pp. 39–45). Florianopolis, Brazil: IEEE. doi: 10.1109/SBESC.2011.18
- Ehrig, H., Ehrig, K., Prange, U., & Taentzer, G. (2006). *Fundamentals of algebraic graph transformation (monographs in theoretical computer science. An EATCS series)*. Berlin, Heidelberg: Springer-Verlag.
- Eker, J., Janneck, J., Lee, E., Jie Liu, Xiaojun Liu, Ludvig, J., ... Yuhong Xiong (2003, January). Taming heterogeneity - the Ptolemy approach. *Proceedings of the IEEE*, 91(1), 127–144. doi: 10.1109/JPROC.2002.805829
- Eker, S., Meseguer, J., & Sridharanarayanan, A. (2004, April). The Maude LTL Model Checker. *Electronic Notes in Theoretical Computer Science*, 71, 162–187. doi: 10.1016/S1571-0661(05)82534-4
- Engels, G., Küster, J. M., Heckel, R., & Groenewegen, L. (2001, September). A methodology for specifying and analyzing consistency of object-oriented behavioral models. *ACM SIGSOFT Software Engineering Notes*, 26(5), 186–195. doi: 10.1145/503271.503235
- Farzan, A., & Meseguer, J. (2007, July). Partial Order Reduction for Rewriting Semantics of Programming Languages. *Electronic Notes in Theoretical Computer Science*, 176(4), 61–78. doi: 10.1016/j.entcs.2007.06.008
- France, R., & Rumpe, B. (2007, May). Model-driven Development of Complex Software: A Research Roadmap. In *Future of Software Engineering (FOSE '07)* (pp. 37–54). Minneapolis, MN, USA: IEEE. doi: 10.1109/FOSE.2007.14
- Gadducci, F. (2007). Graph rewriting for the  $\pi$ -calculus. *Mathematical Structures in Computer Science*, 17(3), 407–437. doi: 10.1017/S096012950700610X
- Kienzle, J., Mussbacher, G., Combemale, B., & Deantoni, J. (2019, October). A unifying framework for homogeneous model composition. *Software & Systems Modeling*, 18(5), 3005–3023. doi: 10.1007/s10270-018-00707-8
- Klare, H., & Gleitze, J. (2019, September). Commonalities for Preserving Consistency of Multiple Models. In *2019 ACM/IEEE 22nd International Conference on Model Driven Engineering Languages and Systems Companion (MODELS-C)* (pp. 371–378). Munich, Germany: IEEE. doi: 10.1109/MODELS-C.2019.00058
- Kräuter, T. (2021). Towards behavioral consistency in heterogeneous modeling scenarios. In *2021 ACM/IEEE international conference on model driven engineering languages and systems companion (MODELS-C)*. doi: 10.1109/MODELS-C53483.2021.00107
- Kräuter, T. (2022, October). *Artifacts - Towards behavioral consistency in multimodeling*. <https://github.com/timKraeuter/Towards-behavioral-consistency-in-multi-modeling>.
- Kräuter, T., König, H., Rutle, A., & Lamo, Y. (2022). Formalization and analysis of BPMN using graph transformation systems. In *EasyChair Preprint no. 8626*.
- Küster, J., & Stehr, J. (2003). Towards explicit behavioral consistency concepts in the UML. In *Proceedings of 2nd ICSE workshop on scenarios and state machines: Models, algorithms, and tools (portland, USA)*.
- Object Management Group. (2013, December). *Business Process Model and Notation (BPMN), Version 2.0.2*. <https://www.omg.org/spec/BPMN/>.
- Object Management Group. (2017, December). *Unified Modeling Language, Version 2.5.1*. <https://www.omg.org/spec/UML>.
- Peter, D. (2022). *Hyperfine*. GitHub.
- Rensink, A. (2004). The GROOVE simulator: A tool for state space generation. In J. L. Pfaltz, M. Nagl, & B. Böhlen (Eds.), *Applications of graph transformations with industrial relevance* (pp. 479–485). Berlin, Heidelberg: Springer Berlin Heidelberg.
- Spanoudakis, G., & Zisman, A. (2001, December). Inconsistency Management in Software Engineering: Survey and Open Research Issues. In *Handbook of Software Engineering and Knowledge Engineering* (pp. 329–380). World Scientific. doi: 10.1142/9789812389718\_0015
- Stünkel, P., König, H., Lamo, Y., & Rutle, A. (2021, July). Comprehensive Systems: A formal foundation for Multi-Model Consistency Management. *Formal Aspects of Computing*. doi: 10.1007/s00165-021-00555-2
- Tisi, M., Jouault, F., Fraternali, P., Ceri, S., & Bézin, J. (2009). On the Use of Higher-Order Model Transformations. In R. F. Paige, A. Hartman, & A. Rensink (Eds.), *Model Driven Architecture - Foundations and Applications* (Vol. 5562, pp. 18–33). Berlin, Heidelberg: Springer Berlin Heidelberg. doi: 10.1007/978-3-642-02674-4\_3
- Vara Larsen, M. E., DeAntoni, J., Combemale, B., & Mallet, F. (2015, September). A Behavioral Coordination Operator Language (BCoOL). In *2015 ACM/IEEE 18th International Conference on Model Driven Engineering Languages and Systems (MODELS)* (pp. 186–195). Ottawa, ON, Canada: IEEE. doi: 10.1109/MODELS.2015.7338249
- Weidmann, N., Kannan, S., & Anjorin, A. (2021, June). Tolerance in Model-Driven Engineering: A Systematic Literature Review with Model-Driven Tool Support. *arXiv:2106.01063 [cs]*.
- Yao, S., & Shatz, S. (2006, November). Consistency Checking of UML Dynamic Models Based on Petri Net Techniques. In *2006 15th International Conference on Computing* (pp. 289–297). Mexico city, Mexico: IEEE. doi: 10.1109/CIC.2006.32

## About the authors

**Tim Kräuter** is a Ph.D. student at the Western Norway University of Applied Sciences, Bergen, Norway. His primary research is on integrating heterogeneous behavioral models in multi-model-driven software engineering. Previously he worked as a software developer and acquired a master of science in Information Engineering at the University of Applied Sciences, FHDW

Hannover, Germany. You can contact the author at [tkra@hvl.no](mailto:tkra@hvl.no) or visit <https://timkraeuter.com/>.

**Harald König** is a professor for Computer Science at the University of Applied Sciences, FHDW Hannover, Germany, and an Adjunct Professor at the Department of Computer Science, Electrical Engineering and Mathematical Sciences at the Western Norway University of Applied Sciences, Bergen, Norway. Before entering academia, he worked at SAP in Walldorf and received his Ph.D. in pure Mathematics from Leibniz University in Hannover, Germany. You can contact the author at [harald.koenig@fhdw.de](mailto:harald.koenig@fhdw.de).

**Adrian Rutle** is a professor at the Department of Computer science, Electrical engineering and Mathematical sciences at the Western Norway University of Applied Sciences, Bergen, Norway. Rutle's main interest is applying theoretical results from the field of model-driven software engineering to practical domains and has his main expertise in the development of formal modeling frameworks for domain-specific modeling languages, graph-based logic for reasoning about static and dynamic properties of models, and the use of model transformations for the definition of the semantics of modeling languages. His recent research focuses on multilevel modeling, model repair, multi-model consistency management, modeling and simulation for robotics, digital fabrication, smart systems, and applications of machine learning in model-driven software engineering. You can contact the author at [aru@hvl.no](mailto:aru@hvl.no).

**Yngve Lamo** is a professor at the Department of Computer science, Electrical engineering and Mathematical sciences at the Western Norway University of Applied Sciences, Bergen. Lamo holds a Ph.D. in Computer Science from the University of Bergen, Norway. His research interests span from formal foundations of Model-Based Software Engineering to its applications, especially in Health Informatics. He is currently applying MDSE to create Adaptive Software Technologies for mental Health. You can contact the author at [yla@hvl.no](mailto:yla@hvl.no).

**Patrick Stünkel** is currently working as a postdoctoral researcher at Haukeland University Hospital in Bergen, Norway, where he is working on applying process mining, workflow modeling, and optimization techniques in digital pathology. Before that, Patrick did his Ph.D. at Western Norway University of Applied Sciences on the topics of semantic interoperability and consistency management among heterogeneous software models. You can contact the author at [Patrick.Stunkel@helse-bergen.no](mailto:Patrick.Stunkel@helse-bergen.no) or visit <https://past.corrlang.io/>.